

CSCD43 IO Investigation

Brandon & Henry

Introduction




For this assignment, we were asked to explore the effects of IO on the disk, and to look at the speed of reading and writing from the disk using different block sizes. As we created our code and tested it on various machines, we noted some interesting results that surprised us! See below for our results.



Writing to a NVME SSD on OS X

Testing the experiment on a laptop with a PCI-E SSD and writing a file of size 10MB, we can see that the graph approximates a logarithmic curve, and that after a block size of around 1 MB, we see negligible gains in performance. Furthermore, we also see that you need at least 1KB of block size to achieve acceptable (SSD) speeds.


BLOCK SIZE	BYTES PER SECOND
100	41011340
1000	95967448
10000	111259457
100000	113958815
1000000	118635220
2000000	116472740
3000000	115344245
1500000	117753730
2500000	111011201



Writing to a HDD on Linux

Firing up our school computers, we ran our script on their hard disks and the results we found were interesting to say the least. It appears that they were writing faster than our laptops! Furthermore we see that the block size is optimal at around 1MB.


BLOCK SIZE	BYTES PER SECOND
100	41790995
1000	51360805
10000	91943032
100000	132987566
1000000	128524792
2000000	123994098
3000000	109357742
1500000	90883478
2500000	87114085



Writing to a USB on Linux

Finally we plugged a USB drive into the lab machines and ran our script, which resulted in a similar result, with the optimal block size appearing to be around 1.5MB


BLOCK SIZE	BYTES PER SECOND
100	5518182
1000	36907996
10000	62344917
100000	87670080
1000000	89338360
2000000	93440478
3000000	87133820
1500000	88161653
2500000	91480428



Reading from a NVME SSD on OS X

From observing the read rates, the first thing we see is that block size does not significantly impact their rates, however the optimal size is still somewhere between 1MB and 1.5MB.

BLOCK SIZE	BYTES PER SECOND
100	284082839
1000	321864238
10000	318461196
100000	397030214
1000000	379506641
2000000	359880520
3000000	304887344
1500000	380170316
2500000	375629179



Reading from a HDD on Linux

From observing the read rates, the first thing we see is that block size does not significantly impact their rates, however the optimal size is still somewhere between 1MB and 1.5MB. It's interesting to note that the speeds were close to or faster than the SSD, so something seems wrong

BLOCK SIZE	BYTES PER SECOND
100	331191627
1000	374405631
10000	379679550
100000	370041445
1000000	369890882
2000000	353769413
3000000	404858300
1500000	392603353
2500000	375629179

Comparison & Thoughts

It appears that both the read & write speed of our drives were similar (which is not what we expected since they are rated to be much different). This leads to a few questions and interesting points that we have thought about. First of all, it appears that Linux overall had much better performance in comparison to Mac OS, which suggests maybe EXT4 is better than HSF+. Furthermore, it appears as though the function “fwrite()” may not directly flush to disk, and instead may flush to cache instead, leading to the similar performance between a laptop SSD & desktop HDD, or perhaps the Linux computers contain a super fast network drive in RAID. Another explanation could be that “fwrite()” does not wait for the file to be written to disk, so writing to disk and immediately recording the time does not account for the full time taken, leading to inflated numbers. Furthermore, upon further investigation we believe there is some sort of inefficiency in the code or the compiler optimized code, as performing a direct write to disk by performing an operation such as “dd if=/dev/zero of=/dev/null bs=1024” yields much different numbers than ours, with the SSD’s performing much closer to their expected read/write speeds.